

foreign-c is a C foreign function interface (FFI) library for R7RS Schemes. It is portable in the sense that it supports multiple implementations.

[Issue tracker](#)

[Mailing lists](#)

[Jenkins](#)

- [Installation](#)
- [Documentation](#)
 - [Types](#)
 - [Primitives 1](#)
 - c-type-size
 - define-c-library
 - define-c-procedure
 - c-bytevector?
 - c-bytevector-u8-set!
 - c-bytevector-u8-ref
 - c-bytevector-pointer-set!
 - c-bytevector-pointer-ref
 - [Primitives 2](#)
 - define-c-callback
 - [c-bytevector](#)
 - make-c-null
 - c-null?
 - c-free
 - make-c-bytevector
 - call-with-address-of
 - native-endianness
 - c-bytevector-s8-set!
 - c-bytevector-s8-ref
 - c-bytevector-s16-set!
 - c-bytevector-s16-ref
 - c-bytevector-s16-native-set!
 - c-bytevector-s16-native-ref
 - c-bytevector-u16-set!
 - c-bytevector-u16-ref

- c-bytevector-u16-native-set!
- c-bytevector-u16-native-ref
- c-bytevector-s32-set!
- c-bytevector-s32-ref
- c-bytevector-s32-native-set!
- c-bytevector-s32-native-ref
- c-bytevector-u32-set!
- c-bytevector-u32-ref
- c-bytevector-u32-native-set!
- c-bytevector-u32-native-ref
- c-bytevector-s64-set!
- c-bytevector-s64-ref
- c-bytevector-s64-native-set!
- c-bytevector-s64-native-ref
- c-bytevector-u64-set!
- c-bytevector-u64-ref
- c-bytevector-u64-native-set!
- c-bytevector-u64-native-ref
- c-bytevector-sint-set!
- c-bytevector-sint-ref
- c-bytevector-uint-set!
- c-bytevector-uint-ref
- c-bytevector-ieee-single-set!
- c-bytevector-ieee-single-native-set!
- c-bytevector-ieee-single-ref
- c-bytevector-ieee-single-native-ref
- c-bytevector-ieee-double-set!
- c-bytevector-ieee-double-native-set!
- c-bytevector-ieee-double-ref
- c-bytevector-ieee-double-native-ref
- bytevector->c-bytevector
- c-bytevector->bytevector
- string->c-utf8
- c-utf8->string

◦ Environment variables

Implementation support tables

NOTE Implementation missing from this table does not mean it will not be supported. Either the work on it has not started yet or support for missing implementations is so unfinished that they are not listed here.

Required versions:

- Chibi > 0.11
 - At the time only 0.11 is out so build from git
- Chicken >= 5.4.0 < 6
- Gauche >= 0.9.15
 - Does not yet work with snow-chibi install
- Guile >= 3
 - Does not yet work with snow-chibi install
 - Has include bug, might not work on all situations
- Kawa >= 3.11 and Java >= 22
 - Needs arguments to enable FFI
 - -J-add-exports=java.base/jdk.internal.foreign.abi=ALL-UNNAMED
 - -J-add-exports=java.base/jdk.internal.foreign.layout=ALL-UNNAMED
 - -J-add-exports=java.base/jdk.internal.foreign=ALL-UNNAMED
 - -J-enable-native-access=ALL-UNNAMED
 - -J-enable-preview
 - All needed arguments on one line for copy pasting
 - -J-add-exports=java.base/jdk.internal.foreign.abi=ALL-UNNAMED -J-add-exports=java.base/jdk.internal.foreign.layout=ALL-UNNAMED -J-add-exports=java.base/jdk.internal.foreign=ALL-UNNAMED -J-enable-native-access=ALL-UNNAMED -J-enable-preview
 - So that snow-chibi installed library is found
 - -Dkawa.import.path=/usr/local/share/kawa
 - -Dkawa.import.path=/usr/local/share/kawa/lib
- Mosh >= 0.2.9-rc1
- Racket >= 8.16 [cs]
- Sagittarius >= 0.9.13

- STklos > 2.10
 - At the time only 2.10 is out so build from git

Primitives 1 table

	c- type- size	c- bytevector- u8-set!	c- bytevector- u8-ref	define- c- library	c- bytevector?	define- proced
Chibi	X	X	X	X	X	X
Chicken	X	X	X	X	X	X
Gauche	X	X	X	X	X	X
Guile	X	X	X	X	X	X
Kawa	X	X	X	X	X	X
Mosh	X	X	X	X	X	X
Racket	X	X	X	X	X	X
Sagittarius	X	X	X	X	X	X
STklos	X	X	X	X	X	X
Ypsilon	X	X	X	X	X	X

Primitives 2 table

	define-c-callback
Chibi	
Chicken	X
Gauche	
Guile	X
Kawa	
Mosh	X
Racket	X
Saggittarius	X
STklos	
Ypsilon	X

Test files pass

	primitives.scm	addressof.scm	callback.scm
Chibi	X	X	
Chicken	X	X	X
Gauche	X	X	
Guile	X	X	X
Kawa	X	X	
Mosh	X	X	
Racket	X		
Saggittarius	X	X	X
STklos	X	X	
Ypsilon	X	X	

Installation

Snow-fort

Not yet installable with snow-fort:

- Gauche
 - Use manual installation

<https://snow-fort.org/>

snow-chibi -impls=IMPLEMENTATION install “(foreign c)”

You can test that library is found by your implementation like this:

```
cp tests/hello.scm /tmp/hello.scm
cd /tmp
IMPLEMENTATION hello.scm
```

Manual system wide

Either download the latest release from <https://git.sr.ht/~retropikzel/foreign-c/refs> or git clone, tag, and copy the *foreign* directory to your library directory.

Example installation for Gauche:

```
make SCHEME=gauche
make SCHEME=gauche install
```

With most implementations the make command does not compile anything. When that is the case it will say “Nothing to build on SCHEME_IMPLEMENTATION_NAME.” ### Manual for project

Either download the latest release from <https://git.sr.ht/~retropikzel/foreign-c/refs> or git clone, tag, and copy the *foreign* directory to your library directory.

Example assuming libraries in directory *snow*:

```
git clone https://git.sr.ht/~retropikzel/foreign-c --branch
LATEST_VERSION
cd foreign-c
make SCHEME_IMPLEMENTATION_NAME
cd ..
mkdir -p snow
cp -r foreign-c/foreign snow/
```

With most implementations the make command does not compile anything. When that is the case it will say “Nothing to build on SCHEME_IMPLEMENTATION_NAME.”

Documentation

Types

Types are given as symbols, for example 'int8 or 'pointer.

- int8
- uint8
- int16
- uint16
- int32
- uint32
- int64
- uint64

- char
- unsigned-char
- short
- unsigned-short
- int
- unsigned-int
- long
- unsigned-long
- float
- double
- pointer
 - c-bytevector on Scheme side
- callback
 - Callback function
- void
 - Can not be argument type, only return type

Primitives 1

(c-type-size type)

Returns the size of given C type.

(define-c-library scheme-name headers object-name options)

Takes a scheme-name to bind the library to, list of C headers as strings, shared-object name and options.

The C header strings should not contain "<" or ">", they are added automatically.

The name of the shared object should not contain suffix like .so or .dll. Nor should it contain any prefix like "lib".

Options:

- additional-versions
 - Search for additional versions of shared object, given shared object "c" and additional versions "6" "7" on linux the files "libc", "libc.6", "libc.7" are searched for.
 - Can be either numbers or strings

- additional-paths
 - Give additional paths to search shared objects from

Example:

```
(define-c-library libc
  (list "stdlib.h")
  libc-name
  '((additional-versions (" " "0" "6"))
    (additional-paths ("."))))
```

Notes

- Do not cond-expand inside the arguments, that might lead to problems on some implementations.
- Do not store options in variables, that might lead to problems on some implementations.
- Pass the headers using quote
 - As '(...) and not (list...)
- Pass the options using quote
 - As '(...) and not (list...)

(define-c-procedure *scheme-name shared-object c-name return-type argument-type*)

Takes a scheme-name to bind the C procedure to, shared-object where the function is looked from, c-name of the function as symbol, return-type and argument-types.

Defines a new foreign function to be used from Scheme code.

Example:

```
(define-c-library libc '("stdlib.h") libc-name '("6"))
(define-c-procedure c-puts libc 'puts 'int '(pointer))
(c-puts "Message brought to you by foreign-c!")
```

Notes

- Pass the return-types using quote
 - As '(...) and not (list...)

(c-bytevector? *obj*)

Returns **#t** if *obj* is c-bytevector, otherwise returns **#f**.

(c-bytevector-u8-set! c-bytevector k byte)

If K is not a valid index of c-bytevector the behaviour is undefined.

Stores the byte in element k of c-bytevector.

(c-bytevector-u8-ref c-bytevector k)

If K is not a valid index of c-bytevector the behaviour is undefined.

Returns the byte at index k of c-bytevector.

(c-bytevector-pointer-set! c-bytevector k pointer)

If K is not a valid index of c-bytevector the behaviour is undefined.

Stores the pointer(which is also c-bytevector) in element k of c-bytevector.

(c-bytevector-pointer-ref c-bytevector k pointer)

If K is not a valid index of c-bytevector the behaviour is undefined.

Returns the pointer(which is also c-bytevector) at index k of c-bytevector.

Primitives 2

(define-c-callback scheme-name return-type argument-types procedure)

Takes scheme-name to bind the Scheme procedure to, return-type, argument-types and procedure as in place lambda.

Defines a new Sceme function to be used as callback to C code.

Example:

```
; Load the shared library
(define-c-library libc-stdlib '("stdlib.h") libc-name '("") "6")

; Define C function that takes a callback
(define-c-procedure qsort libc-stdlib 'qsort 'void '(pointer int
```

```

int callback))

; Define our callback
(define-c-callback compare
  'int
  '(pointer pointer)
  (lambda (pointer-a pointer-b)
    (let ((a (c-bytevector-sint-get pointer-a
      (native-endianness) 0))
          (b (c-bytevector-sint-get pointer-b
      (native-endianness) 0)))
      (cond ((> a b) 1)
            ((= a b) 0)
            ((< a b) -1))))))

; Create new array of ints to be sorted
(define array (make-c-bytevector (* (c-type-size 'int) 3)))
(c-bytevector-s32-native-set! array (* (c-type-size 'int) 0) 3)
(c-bytevector-s32-native-set! array (* (c-type-size 'int) 1) 2)
(c-bytevector-s32-native-set! array (* (c-type-size 'int) 2) 1)

(display array)
(newline)
;> (3 2 1)

; Sort the array
(qsort array 3 (c-type-size 'int) compare)

(display array)
(newline)
;> (1 2 3)

```

c-bytevector

Foreign-c c-bytevector interface is copied from R6RS bytevectors, with some added functionality for C null pointers and manual memory management.

(make-c-null)

Returns a null C pointer.

(c-null? *obj*)

Returns **#t** if *obj* is a null C pointer, otherwise returns **#f**.

(c-free *c-bytevector*)

Frees *c-bytevector* from memory.

(call-with-address-of *c-bytevector thunk*)

Calls *thunk* with address pointer of *c-bytevector*.

Since the support for calling C functions taking pointer address arguments, ones prefixed with & in C, varies, some additional ceremony is needed on the Scheme side.

Example:

Calling from C:

```
//void func(int** i);  
func(&i);
```

Calling from Scheme:

```
(define cbv (make-bytevector (c-type-size 'int)))  
(call-with-address-of  
  cbv  
  (lambda (address)  
    (func address)))  
; Use cbv here
```

The passed c-bytevector, in example named cbv, should only be used **after** call to call-with-address-of ends.

(bytevector->c-bytevector *bytevector*)

Returns a newly allocated c-bytevector of the bytes of *bytevector*.

(c-bytevector->bytevector)

Returns a newly allocated bytevector of the bytes of *c-bytevector*.

(native-endianness)

Returns the endianness symbol associated implementation's preferred endianness (usually that of the underlying machine architecture). This may be any <endianness symbol>, including a symbol other than big and little.

(make-c-bytevector *k*)

(make-c-bytevector *k fill*)

Returns a newly allocated c-bytevector of *k* bytes.

If the *fill* argument is missing, the initial contents of the returned c-bytevector are unspecified.

If the *fill* argument is present, it's value must confine to C uint8_t values , it specifies the initial value for the bytes of the c-bytevector

(c-bytevector-s8-set! *c-bytevector k byte*)

If *k* is not a valid index of c-bytevector the behaviour is undefined.

Stores the *byte* in element *k* of *c-bytevector*.

(c-bytevector-s8-ref *c-bytevector k*)

If *k* is not a valid index of c-bytevector the behaviour is undefined.

Returns the byte at index *k* of *c-bytevector*.

(c-bytevector-char-set! *c-bytevector k char*)

If *k* is not a valid index of c-bytevector the behaviour is undefined.

Stores the *char* in element *k* of *c-bytevector*.

(c-bytevector-char-ref *c-bytevector k*)

If *k* is not a valid index of c-bytevector the behaviour is undefined.

Returns the char at index *k* of *c-bytevector*.

(c-bytevector-uchar-set! *c-bytevector k char*)

If *k* is not a valid index of c-bytevector the behaviour is undefined.

Stores the unsigned *char* in element *k* of *c-bytevector*.

(c-bytevector-uchar-ref *c-bytevector k*)

If *k* is not a valid index of *c-bytevector* the behaviour is undefined.

Returns the unsigned char at index *k* of *c-bytevector*.

(c-bytevector-uint-ref *c-bytevector k endianness size*)

(c-bytevector-sint-ref *c-bytevector k endianness size*)

(c-bytevector-uint-set! *c-bytevector k n endianness size*)

(c-bytevector-sint-set! *c-bytevector k n endianness size*)

Size must be a positive exact integer object. If *k*, ..., *k + size - 1* is not valid indices of *c-bytevector* the behavior is unspecified.

The *c-bytevector-uint-ref* procedure retrieves the exact integer object corresponding to the unsigned representation of size *size* and specified by *endianness* at indices *k*, ..., *k + size - 1*.

The *c-bytevector-sint-ref* procedure retrieves the exact integer object corresponding to the two's-complement representation of size *size* and specified by *endianness* at indices *k*, ..., *k + size - 1*. For *c-bytevector-uint-set!*, *n* must be an exact integer object in the interval $\{0, \dots, 256^{\text{size}} - 1\}$.

The *c-bytevector-uint-set!* procedure stores the unsigned representation of size *size* and specified by *endianness* into *c-bytevector* at indices *k*, ..., *k + size - 1*.

The ...-set! procedures return unspecified values.

Examples:

```
(define cbv (make-c-bytevector (c-type-size 'int)))
(c-bytevector-sint-set! cbv 0 100 (native-endianness) (c-type-size
'int))
(c-bytevector-sint-ref cbv 0 (native-endianness) (c-type-size
'int))
> 100
```

(c-bytevector-u16-ref *c-bytevector k endianness*)

(c-bytevector-s16-ref *c-bytevector k endianness*)

(c-bytevector-u16-native-ref *c-bytevector k*)

(c-bytevector-s16-native-ref *c-bytevector k*)

(c-bytevector-u16-set! *c-bytevector k n endianness*)

(c-bytevector-s16-set! *c-bytevector k n endianness*)

(c-bytevector-u16-native-set! *c-bytevector k n*)

(c-bytevector-s16-native-set! *c-bytevector k n*)

K must be a valid index of *c-bytevector* ; so must $k + 1$. For *c-bytevector-u16-set!* and *c-bytevector-u16-native-set!*, *n* must be an exact integer object in the interval $\{0, \dots, 2^{16} - 1\}$. For *c-bytevector-s16-set!* and *c-bytevector-s16-native-set!*, *n* must be an exact integer object in the interval $\{-2^{15}, \dots, 2^{15} - 1\}$.

These retrieve and set two-byte representations of numbers at indices *k* and $k + 1$, according to the endianness specified by *endianness*. The procedures with *u16* in their names deal with the unsigned representation; those with *s16* in their names deal with the two's-complement representation.

The procedures with *native* in their names employ the native endianness, and work only at aligned indices: *k* must be a multiple of 2.

The ...-set! procedures return unspecified values.

(c-bytevector-u32-ref *c-bytevector k endianness*)

(c-bytevector-s32-ref *c-bytevector k endianness*)

(c-bytevector-u32-native-ref *c-bytevector k*)

(c-bytevector-s32-native-ref *c-bytevector k*)

(c-bytevector-u32-set! *c-bytevector k n endianness*)

(c-bytevector-s32-set! *c-bytevector k n endianness*)

(c-bytevector-u32-native-set! *c-bytevector k n*)

(c-bytevector-s32-native-set! *c-bytevector k n*)

$K, \dots, k + 3$ must be valid indices of bytevector. For `c-bytevector-u32-set!` and `bytevector-u32-native-set!`, n must be an exact integer object in the interval $\{0, \dots, 2^{32} - 1\}$. For `bytevector-s32-set!` and `bytevector-s32-native-set!`, n must be an exact integer object in the interval $\{-2^{31}, \dots, 2^{32} - 1\}$.

These retrieve and set four-byte representations of numbers at indices $k, \dots, k + 3$, according to the endianness specified by *endianness*. The procedures with `u32` in their names deal with the unsigned representation; those with `s32` with the two's-complement representation.

The procedures with `native` in their names employ the native endianness, and work only at aligned indices: k must be a multiple of 4.

The `...-set!` procedures return unspecified values.

(c-bytevector-u64-ref *c-bytevector k endianness*)

(c-bytevector-s64-ref *c-bytevector k endianness*)

(c-bytevector-u64-native-ref *c-bytevector k*)

(c-bytevector-s64-native-ref *c-bytevector k*)

(c-bytevector-u64-set! *c-bytevector k n endianness*)

(c-bytevector-s64-set! *c-bytevector k n endianness*)

(c-bytevector-u64-native-set! *c-bytevector k n*)

(c-bytevector-s64-native-set! *c-bytevector k n*)

$K, \dots, k + 7$ must be valid indices of *c-bytevector*. For *c-bytevector-u64-set!* and *c-bytevector-u64-native-set!*, n must be an exact integer object in the interval $\{0, \dots, 264 - 1\}$. For *c-bytevector-s64-set!* and *c-bytevector-s64-native-set!*, n must be an exact integer object in the interval $\{-263, \dots, 264 - 1\}$.

These retrieve and set eight-byte representations of numbers at indices $k, \dots, k + 7$, according to the endianness specified by *endianness*. The procedures with *u64* in their names deal with the unsigned representation; those with *s64* with the two's-complement representation.

The procedures with *native* in their names employ the native endianness, and work only at aligned indices: k must be a multiple of 8.

The *...-set!* procedures return unspecified values.

(c-bytevector-ieee-single-native-ref)

(c-bytevector-ieee-single-ref)

$K, \dots, k + 3$ must be valid indices of *c-bytevector*. For *c-bytevector-ieee-single-native-ref*, k must be a multiple of 4.

These procedures return the inexact real number object that best represents the IEEE-754 single-precision number represented by the four bytes beginning at index k .

(c-bytevector-ieee-double-native-ref)

(c-bytevector-ieee-double-ref)

$K, \dots, k + 7$ must be valid indices of *c-bytevector*. For *c-bytevector-ieee-double-native-ref*, k must be a multiple of 8.

These procedures return the inexact real number object that best represents the IEEE-754 double-precision number represented by the eight bytes beginning at index k .

(c-bytevector-ieee-single-native-set!)

(c-bytevector-ieee-single-set!)

$K, \dots, k + 3$ must be valid indices of *c-bytevector*. For *c-bytevector-ieee-single-native-set!*, k must be a multiple of 4.

These procedures store an IEEE-754 single-precision representation of x into elements k through $k + 3$ of *bytevector*, and return unspecified values.

(c-bytevector-ieee-double-native-set!)

(c-bytevector-ieee-double-set!)

$K, \dots, k + 7$ must be valid indices of *bytevector*. For *c-bytevector-ieee-double-native-set!*, k must be a multiple of 8.

These procedures store an IEEE-754 double-precision representation of x into elements k through $k + 7$ of *bytevector*, and return unspecified values.

(string->c-utf8 *string*)

Returns a newly allocated (unless empty) *c-bytevector* that contains the UTF-8 encoding of the given string.

(c-utf8->string *c-bytevector*)

Returns a newly allocated (unless empty) string whose character sequence is encoded by the given *c-bytevector*.

Utilities

libc-name

Name of the C standard library on the current operating system.
Supported OS:

- Windows
- Linux
- Haiku

See `foreign/c/libc.scm` to see which headers are included and what shared libraries are loaded.

Example:

```
(define-c-library libc '("stdlib.h") libc-name '(" " "6"))  
(define-c-procedure c-puts libc 'puts 'int '(pointer))  
(c-puts "Message brought to you by foreign-c!")
```

Environment variables

Setting environment variables like this on Windows works for this library:

```
set "FOREIGN_C_LOAD_PATH=C:\Program Files (x86)/foo/bar"
```

FOREIGN_C_LOAD_PATH

To add more paths to where foreign c looks for libraries set `FOREIGN_C_LOAD_PATH` to paths separated by `;` on windows, and `:` on other operating systems.